BUILD WHAT'S NEXT: AI FOR PMS

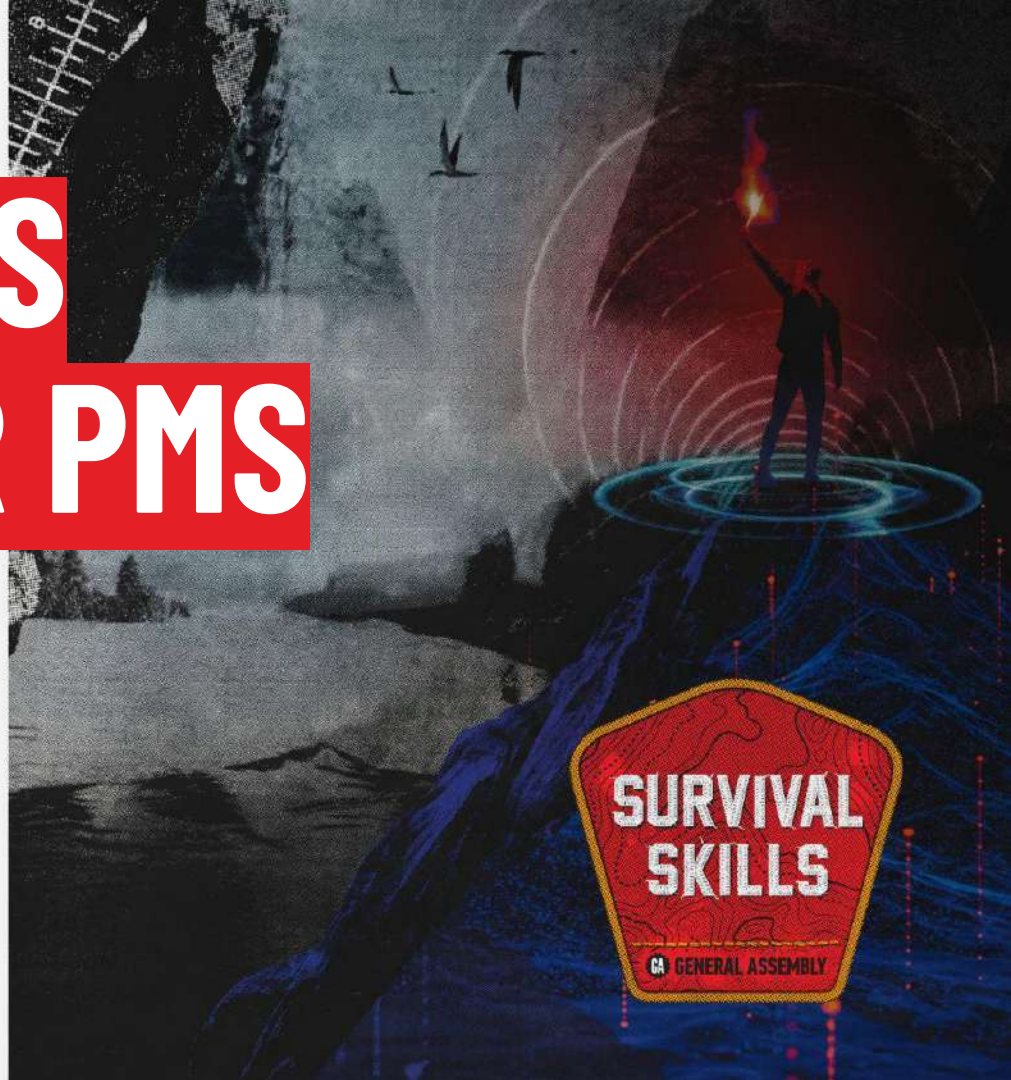

SKILLS ARE YOUR SURVIVAL - FREE CLASS SERIES

Course Roadmap

- Warm up + Welcoming
- On The Agenda
- Generative AI impact and LLM
- Leveraging AI
- Q&A
- Survey+ Closing

General Assembly

General Assembly is at the **forefront of tech talent development**, specialising in today's most in-demand skills.

We **train and connect diverse tech talent** directly to employers looking to **fill their talent needs.**

CONTENT AREAS

- Engineering & Development
- Data & AI
- Product & UX
- Digital Marketing

FAST FACTS

- **110K+** global alumni
- **25,000+** employees of Fortune 500 companies trained in **20+** countries
- **Flagship campuses** in New York & Singapore + global virtual learning in **over 50** markets
- A brand of the **Adecco Group**, the world's foremost provider of staffing, career transition and talent development solutions.

US Financing Options

Options vary by region. Talk to admissions to see what you qualify for.

Installments Divide tuition into three easy payments.

Loans Apply for a loan from one of our high-quality lenders.

Full Tuition Save on tuition by paying in full upfront.

Scholarships Our admissions team will help you determine if you're eligible for a scholarship.

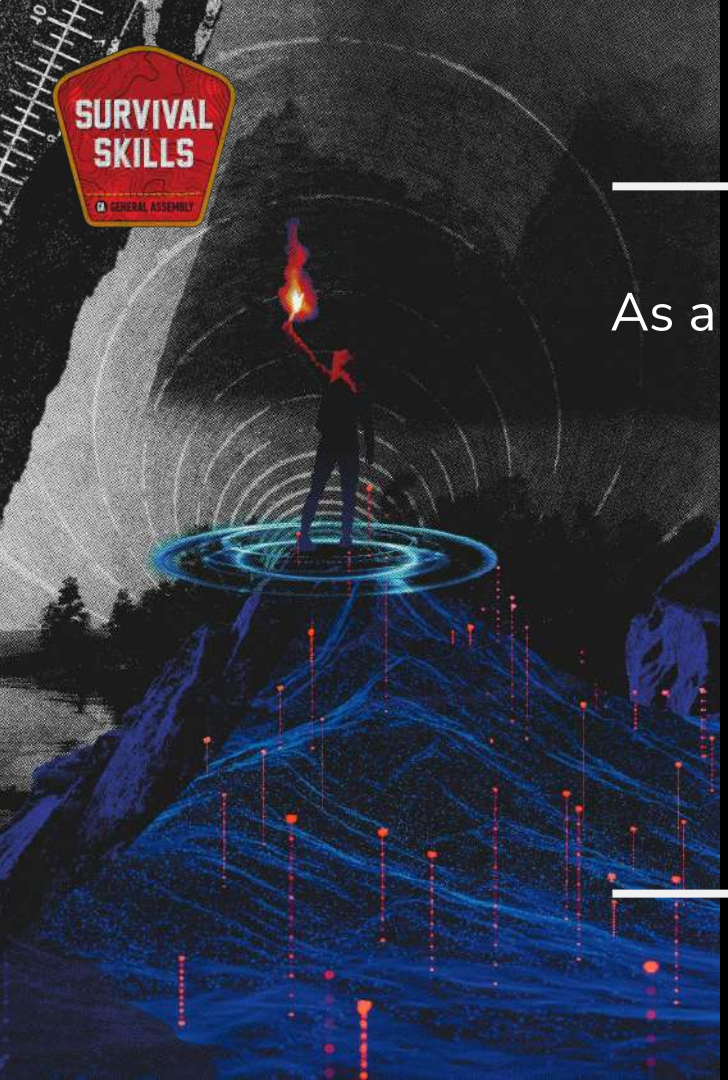

As a thank you for joining us today, we're offering:

$25 off a workshop with code Survive26

+ No application fee for new AI courses

Enroll by May 15, 2026

Terms & conditions apply. Discount amount varies based on region and course selected.
Mention Survival Skills when you speak to Admissions to redeem for a course.

Before we get started...

Some quick reminders

- This is a beginner-level overview of Build What's Next: AI for PMs.
- There will be time for live Q&A at the end.
- YES—we're recording. You'll get the recording from the class via email within 3-5 business days.

How we'll use the Zoom classroom

Set your chat box to "Everyone" so your fellow classmates can see what you have to say.

Use the Q&A as a "parking lot" for general subject matter questions.

Use reactions to make it extra fun and interactive.

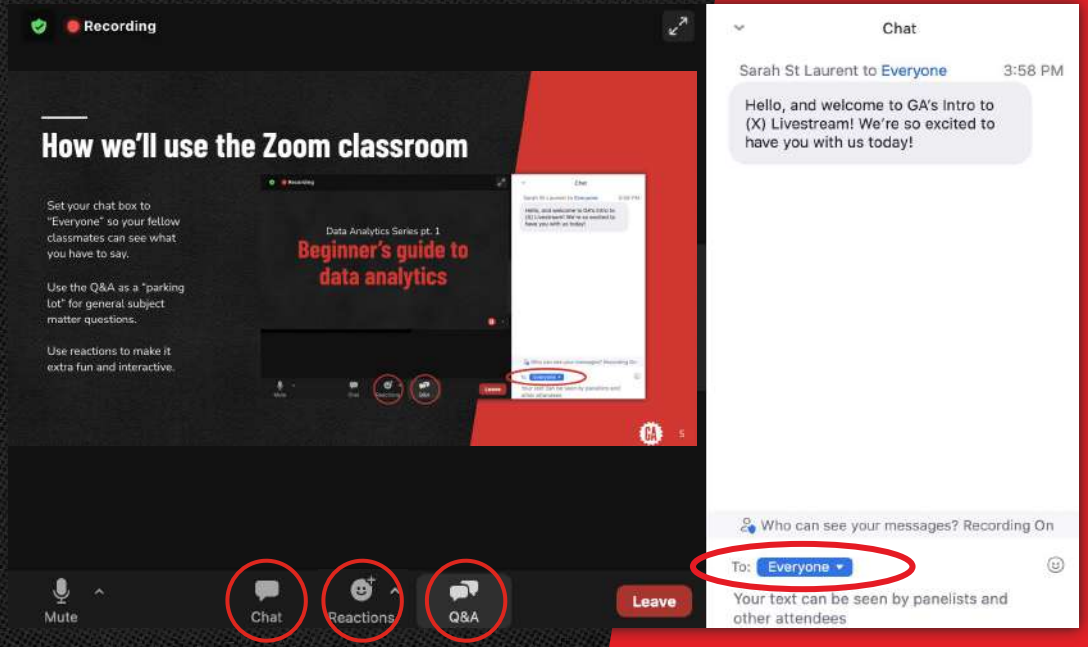

Now that we got that out of the way, let's dive in.

SURVIVAL SKILLS
GENERAL ASSEMBLY

On the Agenda...

- Large Language Models, explained
- Apply AI techniques to improve clarity and efficiency in business writing and communication
- Leverage AI to take on admin tasks
- Use AI to streamline and enhance a multi-step decision-making process

Large Language Models Explained.

AI is a Broad Category

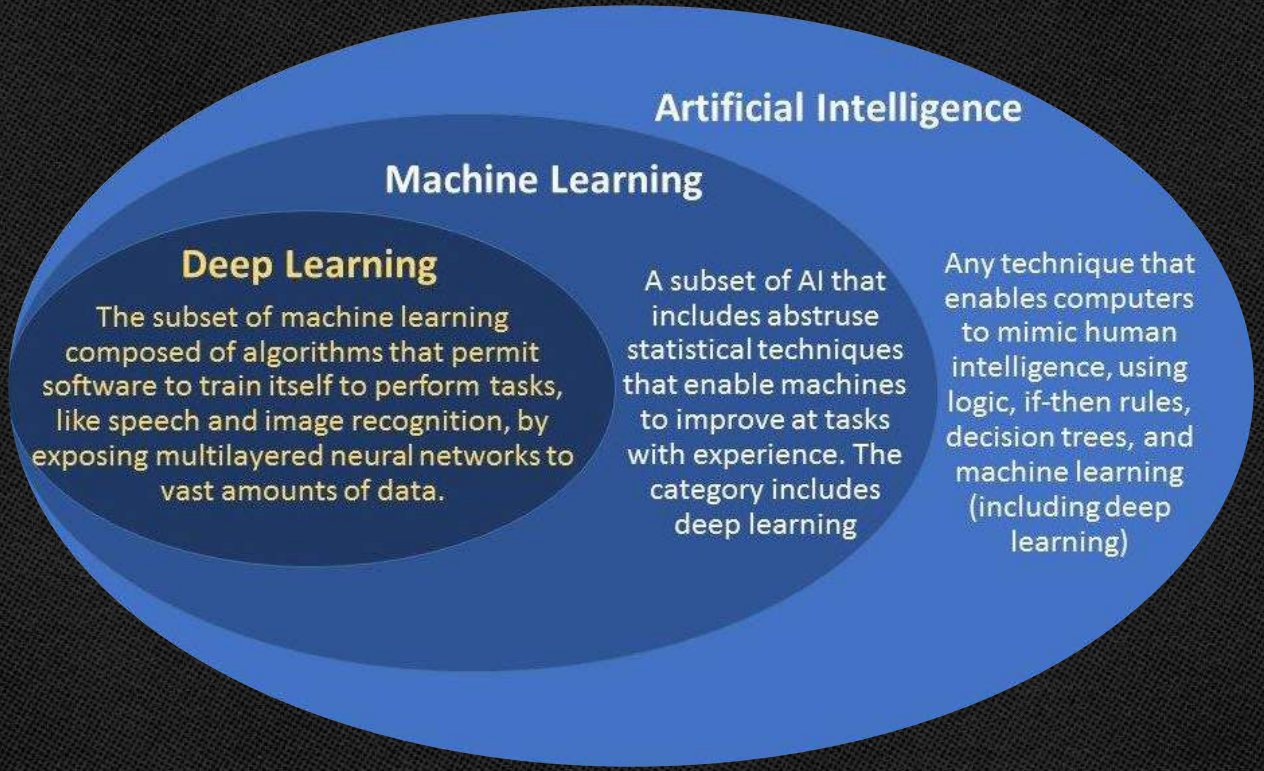

IBM WATSON

What is Machine Learning?

Math and statistics that allow you to **predict trends from historical data**

Analytics can show what *did* happen; **machine learning** allows you to mathematically estimate what *might* happen.

What can Machine Learning Determine?

1. **Is this A or B?** ➡ Classification
2. **Is this weird?** ➡ Anomaly detection
3. **How much/how many?** ➡ Regression analysis
4. **How is it organized?** ➡ Unsupervised learning (*e.g., clustering)*
5. **What should I do next?** ➡ Reinforcement learning

3 Major Machine Learning Categories

Banking Example:

Supervised Learning

Predict credit worthiness of credit card holders:

Build a machine learning model to look for delinquency attributes by providing it with data on delinquent and non-delinquent customers.

Unsupervised Learning

Segment customers by behavioral characteristics:

Survey prospects and customers to develop multiple segments using clustering.

Reinforcement Learning

Create a 'next best offer' model for the call center group:

Build a predictive model that learns over time as users accept or reject offers made by the sales staff.

Supervised Learning

Supervised learning is done with a ground truth, giving the algorithm prior knowledge of what the output values for the data should be.

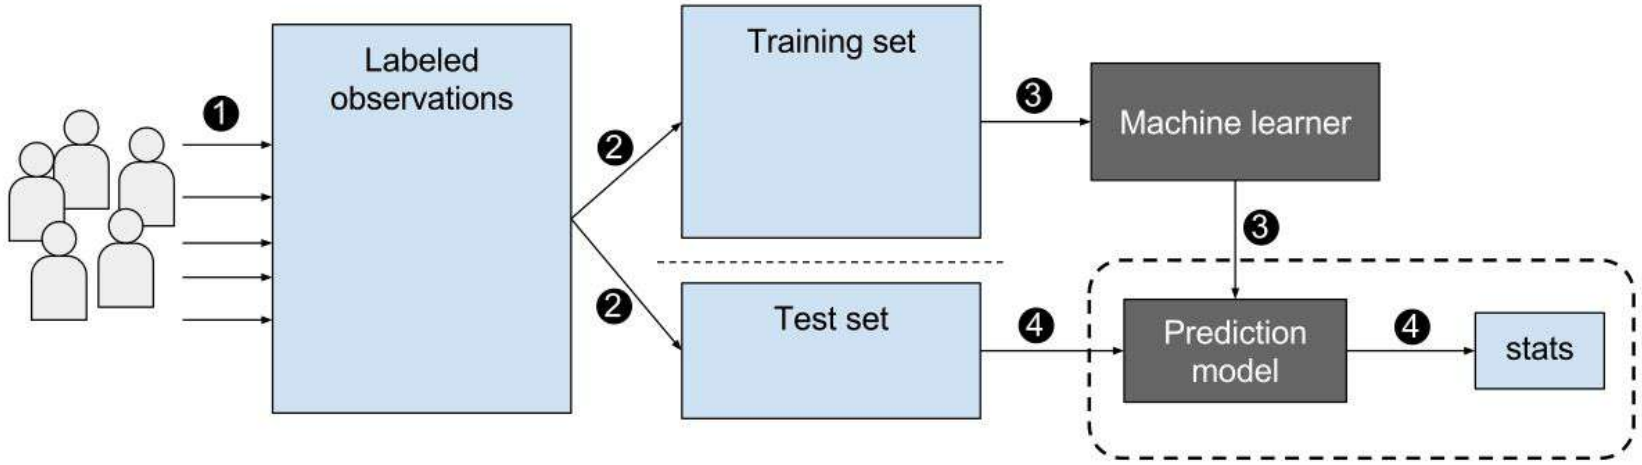

Popular Tools

Platform	System	Models	Features	Cost
OpenAI Backed by Microsoft	ChatGPT	GPT 5	Fast, cheap	Free
	ChatGPT Plus /Enterprise	GPT 4-V/o DALL•E 3	Code, web, multimodal	$20/month
	Playground	Multiple	Tweak Parameters	$5+ PAYG
Microsoft Powered by OpenAI	CoPilot, Bing AI *(Creative Mode)*	GPT 4-V DALL•E 3	Internet, Multimodal	Free
Anthropic Backed by Amazon	Claude AI	Sonnet 3.7	Large Documents	Free or $20/month
Google	Gemini	Gemini	Connect to Gmail, Drive	Varies

Meeting Summaries - Take Caution ⚠️

Samsung used AI to generate meeting minutes. This resulted in a leak of proprietary information about their semiconductors!

Whoops, Samsung workers accidentally leaked trade secrets via ChatGPT

ChatGPT doesn't keep secrets.

By Cecily Mauran on April 6, 2023

LLMs are designed to be useful for everyone –

PROMPT ENGINEERING

helps the tool work for YOU!

Mastering these skills is not *easy*.

In many ways, it's a game of probability, trial and error, continuous learning.

Types of AI Tools in PM

AUTOMATION

🧠 Predictive

Analyzes patterns to forecast outcomes (e.g., attrition risk)

🛠️ Prescriptive

Recommends actions based on data (e.g., benefits optimization)

AUGMENTATION

✍️ Generative AI

Produces new content like text or images (e.g., ChatGPT)

This course focuses mostly on <u>generative AI</u>

How a Predictive & Prescriptive (aka Discriminative AI) Work

Pattern Recognition in Data

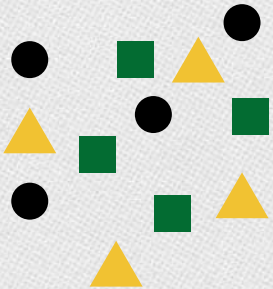

Classification / Categorization

Circles are black; triangles are yellow; squares are purple.

How a Generative AI Works

Pattern Recognition in Data

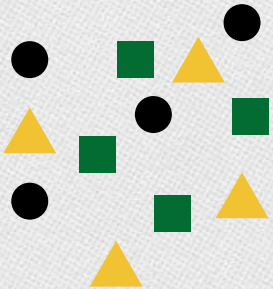

Pattern Recognition in Sequence

Once upon a time, there were three little pigs who set out to build their own houses. The first little pig built his house out of straw. The second little pig built his house out of sticks. The third little pig built his house out of bricks.

Classification / Categorization

Circles are black; triangles are yellow; squares are purple.

Probability & Relationships of Sequence

If the word “house” appears, then “out of” is likely to be the next two words after it.

What AI Can Do in Product Management

- ➔ Draft product requirement documents (PRDs) and feature specifications
- ➔ Screen feature requests using custom prioritization prompts
- ➔ Draft go-to-market plans and strategic decision memos
- ➔ Summarize user interview notes or sentiment data
- ➔ Generate first drafts of competitive analysis reports
- ➔ Develop user persona profiles based on research data

What AI Can't or Shouldn't Do in Product Management

- ➔ Make strategic product decisions
- ➔ Build genuine customer empathy
- ➔ Navigate complex organizational politics
- ➔ Innovate breakthrough product concepts
- ➔ Provide leadership when products face challenges in market

ChatGPT: Limitations and Guidelines

Ideation

Task	Is ChatGPT useful?
Generating ideas for your target audience	✓
Choosing what's best for your product	✗
Providing market insights	✓
Providing information on the competitor landscape	✓
Making a go/no-go decision	✗
Creating a strategic roadmap	✗

Implementation

Task	Is ChatGPT useful?
Coming up with the flow and elements of an experience you are building	✓
Doing actual UX work	✗
Generating snippets of the code you'll need for your product	✓
Doing the actual coding for your product	✓
Inspiring and leading your team	✗

Pre-Launch

Task	Is ChatGPT useful?
Launching your product for you	✗
Coming up with a roll-out plan	✓
Creating marketing content for you	✓
Predicting trends for your product	✓

#23WIP

Key takeaways

AI can automate some things in PM, and augment many others

AI tools are powerful partners for PM professionals

AI doesn't remove bias, judgment, or responsibility

Collaborating with AI: From Simple Requests to Powerful Solutions

- Agents
- Custom Bots
- Multi-Context Projects
- Model Selection & Optimization
- Contextual Prompting with Files & LLMs
- Structured Prompt Engineering
- Basic Prompting

Complex

Simple

All models are

WRONG

but some are useful.

George E. P. Box, Statistician

What do we *need* to know about Large Language Models?

Trained on vast amounts of **internet text**—*hundreds of billions of words*

Outputs are **fluent**, **but not always accurate**

Doesn't understand context or meaning like a human

Not connected to live data or internal company systems

They work by **predicting the most likely next word** in a sentence

How ChatGPT Understands and Responds

ChatGPT sees my prompt as tokens. Then to respond, predicts what tokens should come next.

Clear | Show example

Tokens	Characters
19	89

ChatGPT sees my prompt as tokens. Then to respond, predicts what tokens should come next.

https://platform.openai.com/tokenizer

Why Prompting Matters

*This prompt gives more **structure and signal** to the model.*

Less Effective Prompt

"Tell me about marketing."

Tokens	Characters
5	24

Tell me about marketing.

More Effective Prompt

"Act as a CMO. Summarize 3 marketing trends in Gen Z fashion, with examples."

Tokens	Characters
22	77

Act as a CMO. Summarize 3 marketing trends in Gen Z fashion, with examples.

What is the right

MINDSET

to have when working with an LLM?

Think of ChatGPT as:

An assistant who has read everything useful on the internet, but only remembers it vaguely.

➔ It's *fast, confident, and helpful*—but it:

- Can make things up (a.k.a. hallucinate)
- Misses nuance unless prompted carefully
- Benefits from context and follow-up

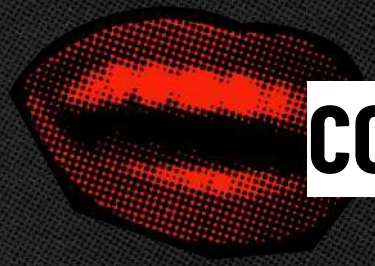

CONSTANTLY EVALUATE OUTPUTS

Research has shown that quality of work can actually DECREASE when someone works with a high-performing model.

Researchers refer to this as “falling asleep at the wheel.”

Source

Never

FALL ASLEEP

at the wheel!

1

High-performing AI **can lower your guard**

2

Users of better AI are **more likely** to accept errors

3

Critical thinking must **increase**, not decrease

4

PMs must **partner** with AI, not hand it the keys

Prompt Creation: specific to PM research

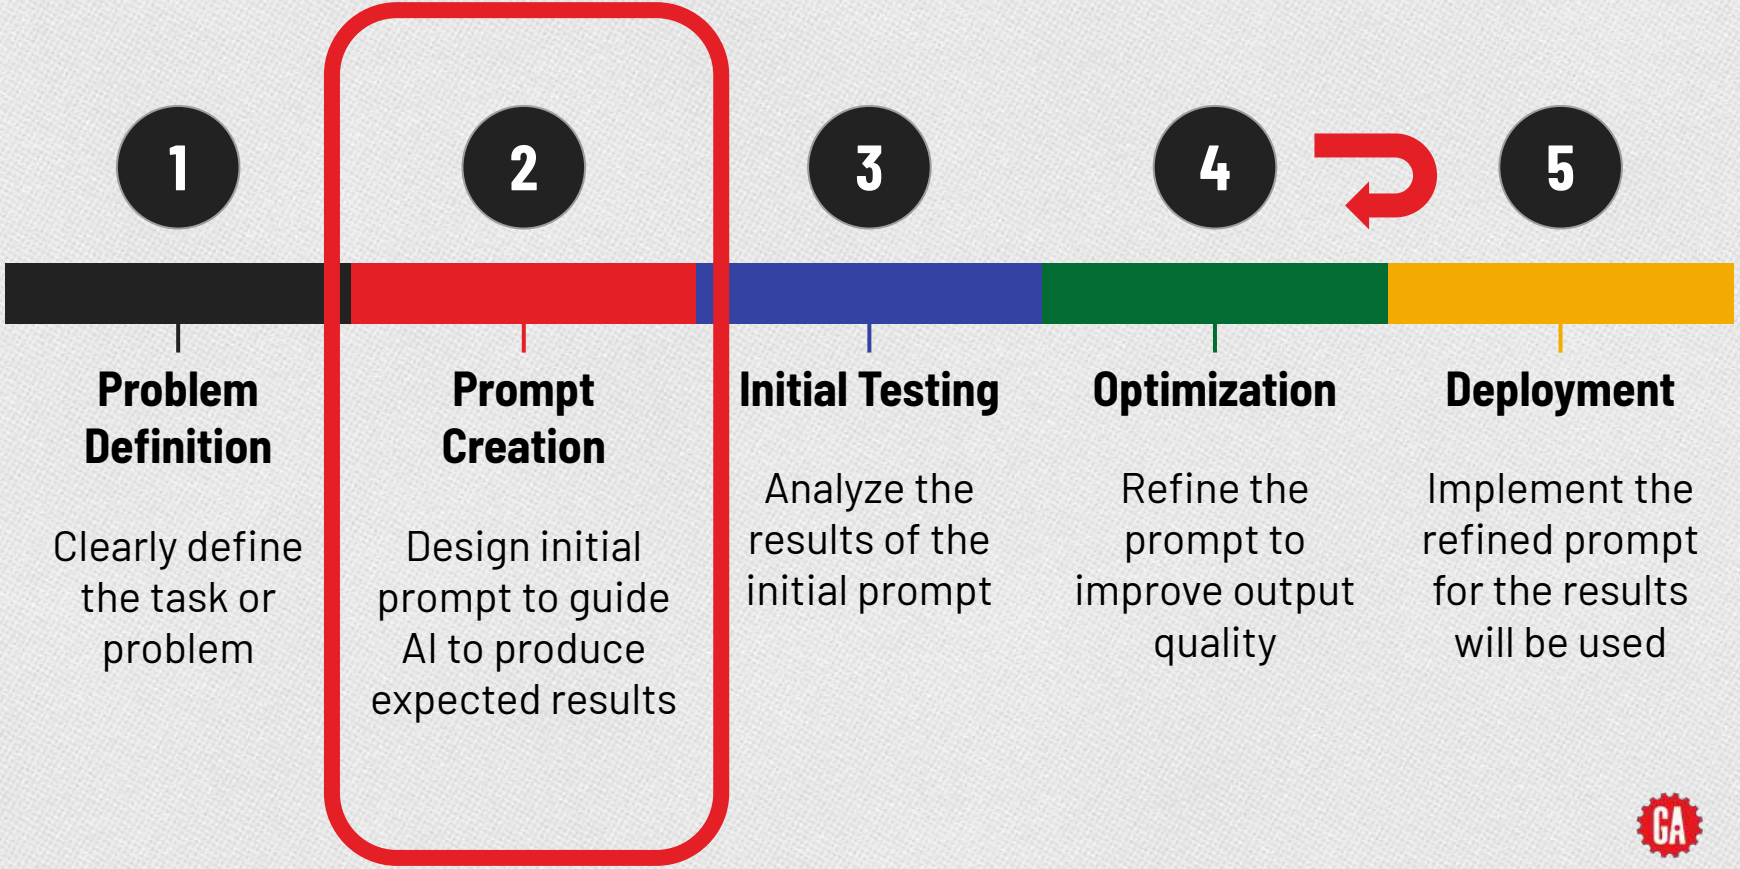

Structuring Research Prompts: TOPS Framework

TASK Definition
What are we research and why?

Define research scope, including setting business context and background.

OUTPUT Specification
What format should the response take?

Provide explicitly formatting requirements. Feel free to use industry jargon or provide examples.

PROCESS guidance
How should the AI go about doing the task?

Provide methodological instructions, including specific steps you want taken.

STANDARDS setting
How will we know it's good?

Establish quality criteria, and *ask* the AI to evaluate the response before sending the results to you.

Structuring Research Prompts: TOPS Framework

TASK
Definition
What are we research and why?

You are a business research analyst. I need background on the company Acme Logistics as I prepare for a potential B2B partnership pitch. Focus on their market presence, recent news, and potential synergies with our company.

OUTPUT
Specification
What format should the response take?

Provide a 3-paragraph summary:

1. *Overview of the company (size, sector, core offerings)*
2. *Summary of recent news and strategic moves*
3. *Potential partnership opportunities.*

Format it as a professional briefing note.

PROCESS
guidance
How should the AI go about doing the task?

Pull data from credible, publicly available sources. Prioritize recent articles and press releases from the last 12 months. If there's a gap in data, explicitly note it.

STANDARDS
setting
How will we know it's good?

Ensure all claims are supported with links or citations. Summarize key data points clearly. Do not include speculation.

Evaluate whether this output meets professional briefing standards, and evaluate before you respond.

Testing & Optimization: specific to PM research

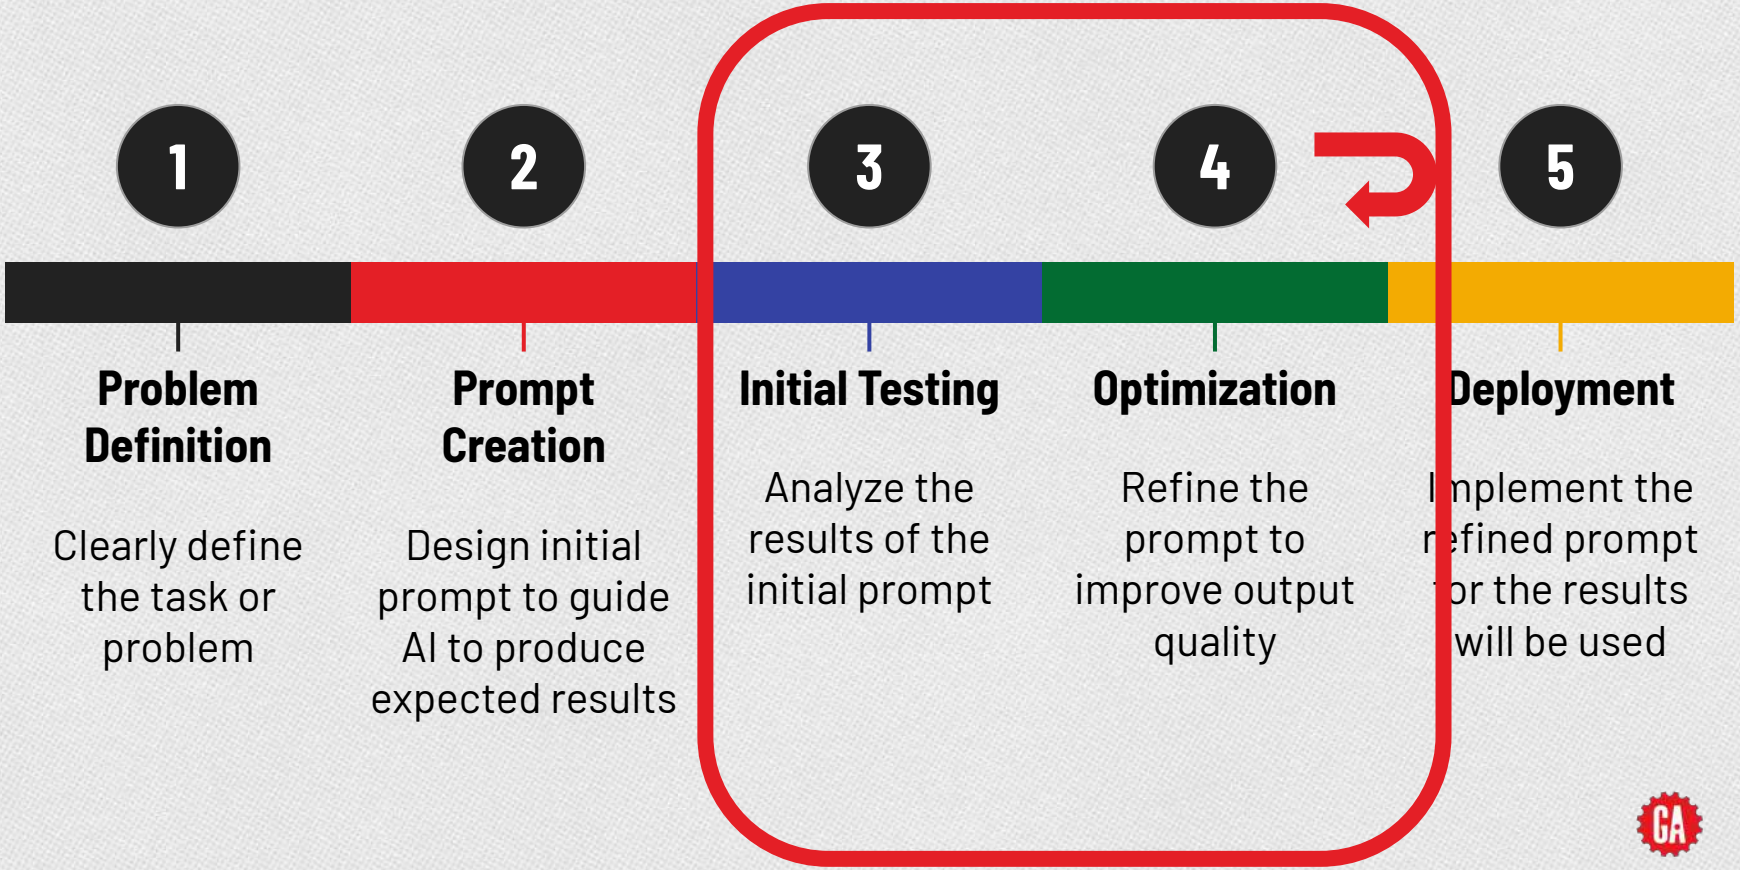

Strategic Levers for Iterating a Prompt

Iteration Strategy	Description	Example
Add missing context	Give background that helps the AI understand the "why."	*"I'm presenting to executive stakeholders who care about profitability."*
Specify format or tone	Shape the response to fit your use case.	*"Make it bullet points with plain language."*
Reference prior output	Build on what's already been generated.	*"Using the list you just gave me, draft a summary for leadership."*
Refine your ask	Be clearer about what you want.	*"Focus more on competitor pricing than general strategy."*
Add constraints	Use limits to guide output.	*"Keep it under 150 words."*
Provide feedback directly	Treat the LLM like a teammate.	*"This was too generic—dig deeper into their AI strategy."*

Iteration Isn't Repetition - It's Strategy

You are the Master. The AI is the student.

Every iteration should build upon learnings from the previous ones.

Iteration allows you to:

- Identify hallucinations
- Fill in missing context
- Narrow or expand the scope
- Clarify ambiguous request
- Add edge cases
- Provide other feedback to the AI

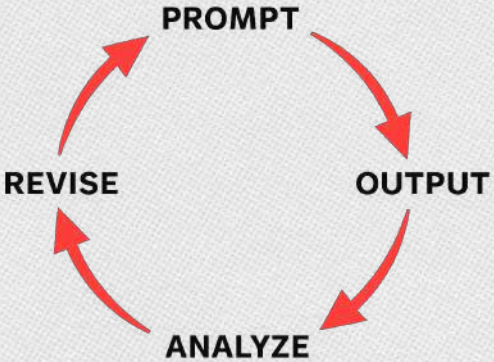

Spotting & Fixing AI Hallucinations

DEFINITION

Content generated by an AI is **factually incorrect, fabricated, or misleading**, even though it might sound plausible.

WHY DO THEY HAPPEN?

- AI tries to **fill gaps** with most probably language (not most accurate)
- AI doesn't have live access to **all** sources
- Your **prompt might be vague** or ask for something **the model wasn't trained to answer**

"95% of market share in 2023..."

What source are you using?

"Oh, you're right. I don't have a source. That must not be true."

Spotting & Fixing AI Hallucinations

How to prevent or fix them in your prompt

Strategy	Example Prompt Tweak
Require Sources	“Cite only from published company press releases or analyst reports.”
Narrow the Scope	“Limit to what Acme has done in the past 12 months.”
Request Uncertainty	“If uncertain, say so rather than guess.”
Use Model Memory Wisely	“Do not invent company names or URLs.”

Three Power Tips for Research Prompting

1. Chunking versus All-at-Once
2. Evaluation criteria for Research Outputs
3. Iteration Patterns

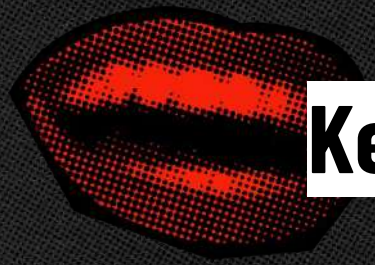

Key takeaways

TOPS

Use right prompt framework to get best results.

Iterate

Have a dialogue with the LLM to dig deeper and find insights.

Deepen Value

Your time is valuable. Use it on things that bring impact. Delegate everything else.

Evaluate LLM's Output

FIVE PILLARS OF ETHICAL AI

EXPLAINABILITY

How and why does it arrive at a particular outcome?

FAIRNESS

Does it treat individual groups equitably?

ROBUSTNESS

Can it handle exceptional conditions?

TRANSPARENCY

Is enough info shared on how it was designed or developed?

PRIVACY

Does it prioritize and safeguard human privacy and data rights?

Ethics of AI in Product Management

⚖️ **Bias & Fairness**	Are you inadvertently excluding demographic segments?
🔍 **Transparency**	Can you explain which parts of an analysis are AI-generated versus human judgement?
🔐 **Privacy**	Is customer data being fed into AI tool?
🛋️ **Accountability**	Who's responsible for what decisions?

AI will make mistakes... what are the consequences when it does?

👏 – A Rubric to Evaluate AI Inputs and Outputs

CLAP Rubric	BEFORE YOU ENTER A PROMPT:	AFTER THE LLM RESPONDS:
Confidentiality	Does this prompt include private, sensitive, or personally identifiable information (PII)?	Does the response include or infer any private or sensitive details that shouldn't be there?
Language and Bias	Is the language neutral and respectful? Could it unintentionally reinforce stereotypes or bias?	Does the output contain biased, harmful, or exclusionary language or assumptions?
Accuracy	Am I asking something that's clear, factual, and verifiable?	Are the facts, references, or assumptions in the response accurate and trustworthy?
Property	Am I entering content that belongs to someone else (e.g., proprietary code, copyrighted text)?	Does the response replicate or reuse proprietary, licensed, or copyrighted material?

AI in Product

AI isn’t magic.

It’s a tool.

Can AI Bring Real Product Value?

How to evaluate product value of AI?

Use frameworks, not hype

Scale your product thinking

Make more informed, faster decisions

Your job isn't to trust it – it's to *evaluate* it.
It's a **resource**.

A Problem for AI? Or... Not?

GOOD FOR AI

- Problem requires p**attern recognition, prediction, generation, or classification**.
- Problem is **repetitive**, data-heavy, time-sensitive.
- Solution could a**utomate, augment, or personalize** a user or team experience.
- Human **decision-making is slow**, subjective, or inconsistent in this area.

BAD FOR AI

- Problem has **low data** availability or data is **highly sensitive**.
- Problem has unclear or **ambiguous** success criteria.
- Solution **requires human nuance** (e.g. conflict resolution)
- **Risk of harm is high** if AI is wrong, biased, or misunderstood.
- Task is one-off or **rare**.

Evaluate AI using

THREE PILLARS

feasibility, viability, desirability

Three-Pillar Evaluation Model

Feasibility

CAN we build it?

Business Viability

SHOULD we build it?

Desirability & Success Metrics

Will customers **WANT** it?

Key takeaways

AI Value

AI needs to provide value; don't do AI for AI's sake.

Business

What are the business risks and feasibility?

Tech

Weigh the technical risk with the feasibility.

Let us know what you thought of today's event:

ga.co/skillsfeedback

View all questions, or only questions you've asked.

Sort by Most Upvotes or Recently Asked

Already see your question posted? Avoid duplicates and **upvote it!**

Send anonymous questions if you're more comfortable. Otherwise, let us know who you are.

You're up.

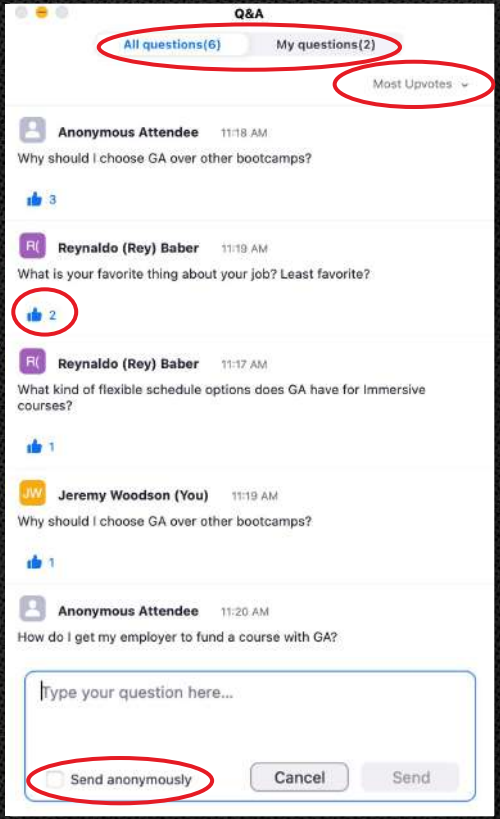

Thank You

GENERAL ASSEMBLY

Can’t wait to see you next time.